

Healio.AI — Senior Developer Critical Analysis

Reviewed: Every directory and ~30 key files across frontend, backend, diagnosis engine, auth, security, and project hygiene.

🔴 SEVERITY: CRITICAL

1. Leaked API Key in Committed `.env` File

[backend/.env](#) contains a **live Resend API key** (`re_NrDRKumJ_...`) committed to Git. This is a **P0 security incident**.

[!CAUTION] **Action Required:** Immediately rotate this key in [Resend dashboard](#), add `backend/.env` to `.gitignore`, and use `git filter-branch` or BFG Repo Cleaner to purge history.

2. `getNotifications()` Bug — `unreadOnly` Filter Silently Ignored

In [api.ts:383-405](#), the function builds a filtered query chain (`query`) but **never executes it**. Instead, it creates a brand-new unfiltered query at line 394. Every call to `getNotifications(userId, true)` returns **all** notifications, not just unread ones.

```
async getNotifications(userId: string, unreadOnly = false) {
    let query = supabase
        .from('notifications')
        .select('*')
        .eq('user_id', userId)
        .order('created_at', { ascending: false });

    if (unreadOnly) {
        query = query.eq('is_read', false);
    }

    - const { data, error } = await supabase
    -     .from('notifications')
    -     .select('*')
    -     .eq('user_id', userId)
    -     .order('created_at', { ascending: false });
    + const { data, error } = await query;
```

3. `createAppointment()` Hardcodes Consultation Fee

In [api.ts:277](#), `consultation_fee: 500` is hardcoded rather than pulled from the doctor's profile. This will **silently overcharge or undercharge** every patient.

4. No Server-Side Route Protection

[rbac.ts](#) provides role-checking utilities, but:

- [middleware.ts](#) only refreshes sessions — it performs **zero role-based access control**.
- A patient can directly navigate to `/admin` or `/doctor` routes by typing the URL.
- All API routes lack server-side auth verification.

[!WARNING] The RBAC system is cosmetic. Any authenticated user can access any portal. The middleware must enforce `isRouteAllowed()` checks after `getUser()`.

● SEVERITY: HIGH — Diagnosis Engine Architecture

5. "Bayesian" Engine Is Not Actually Bayesian

The core engine in [engine.ts](#) claims to use Bayesian inference but implements a **manual log-probability accumulation** with hardcoded boost values (e.g., `+3.0`, `+2.0`, `-5.0`) passed through a sigmoid. This is not Bayesian – it's a **weighted scoring function with a squashing activation**.

What a proper Bayesian network would provide:

Feature	Current Implementation	Proper Bayesian Network / MCMC
Priors	Flat map of 5 prevalence buckets	Learned priors from epidemiological data
Likelihood	Hardcoded boosts (<code>+3.0</code>)	$\text{P}(\text{symptom})$
Posterior	Sigmoid of summed log-scores	True $\text{P}(\text{disease})$
Uncertainty	Heuristic width heuristics	Posterior distribution sampling (MCMC/HMC)
Conditional independence	None — all symptoms treated independently	Directed acyclic graph encoding dependencies
Calibration	None	Platt scaling or isotonic regression
Evidence updates	Re-score from scratch	Incremental belief propagation

[!IMPORTANT] **Recommendation:** Replace the scoring engine with a proper **Bayesian Network** using libraries like `pgmpy` (Python) or `bayes-server`, with:

- A DAG encoding symptom-disease-comorbidity relationships
- MCMC sampling (Metropolis-Hastings or NUTS) for posterior estimation
- Calibrated confidence intervals from posterior distributions
- Network structure learned from clinical data (ICD-10 mappings)

6. Duplicate Scoring Functions

[engine.ts](#) contains **both** `calculateScore()` (lines 228-455) AND `calculateBayesianScore()` (lines 463-731). The former is a dead function — never called. 1,200 lines of engine code that should be ~400.

7. Symptom Matching is Fragile String Matching

All symptom detection relies on `string.includes()` against free-text notes. This produces:

- **False positives:** "I do NOT have chest pain" → matches "chest" + "pain"
- **Missed detections:** "hurts in my torso" won't match "chest"
- The negation regex (`/no|not|without/`) is too simplistic — fails on "I have nothing but chest pain"

Suggestion: Use NLP tokenization + NER (Medical NER like SciSpacy or BioBERT), or at minimum build a proper token-level matching system instead of substring searches.

8. Mock Embedding Fallback Returns Random Vectors

In [embeddings/route.ts](#), when no OpenAI key is available, the API returns a **random 1536-dim vector**. This means vector search in [retrieval.ts](#) returns **random conditions** — the entire diagnosis pipeline degrades to the local keyword fallback silently, with no indication to the user.

9. Retrieval Falls Back to Returning ALL Conditions

In [retrieval.ts:70](#), when vector search yields < 5 results, the system returns `Object.values(CONDITIONS)` — the **entire conditions catalog** — bypassing semantic relevance entirely. This turns the "smart" retrieval into brute force.

🟡 SEVERITY: MEDIUM — Code Quality & Architecture

10. Middleware Uses Deprecated Supabase Cookie Pattern

[middleware.ts](#) uses `get/set/remove` individual cookie methods. The modern `@supabase/ssr` library expects `getAll/setAll` (used correctly in [supabaseServer.ts](#) but not in middleware).

11. 277-Line Mock Data File in Production `lib/`

[mockData.ts](#) contains 8 hardcoded patient records with PII-like data (names, emails, phone numbers, medical conditions). This is imported by production components and should:

- Be moved to `__tests__/fixtures/` or `data/seed/`
 - Never ship in production bundles
-

12. Entire Legacy Vite App Still in Repo

A complete [legacy/](#) directory contains a full Vite application with its own `node_modules/`, `package.json`, and `src/`. This adds **30+ files** and bloats the repo by ~30MB+.

13. Project Root is a Dumping Ground

The root directory contains **20+ loose files** that don't belong:

File Pattern	Count	Action
<code>lint_*.txt, lint_*.log, lint_*.json</code>	8	Delete or add to <code>.gitignore</code>
<code>tsc_*.txt, tsc_*.log</code>	5	Delete
<code>test_*.txt, test_*.sql</code>	6	Move to <code>tests/</code> OR <code>scripts/</code>
<code>build_*.txt, build_*.log</code>	2	Delete
<code>*.sql</code> (loose migrations)	6	Move to <code>supabase/migrations/</code>
<code>app_errors.txt, diagnosis_errors.txt</code>	3	Delete
<code>verification_result.txt</code>	1	Delete

14. Python Backend Has No Authentication

[backend/main.py](#) has rate limiting but **zero authentication**. Any person on the internet can:

- Upload files to the server
 - Send emails through the `/api/email/*` endpoints (spam vector)
 - The email endpoints don't verify the requesting user owns the target email
-

15. Duplicate Dependency in `requirements.txt`

[requirements.txt](#) lists `python-multipart` twice (lines 6 and 8). Dependencies are also unpinned (no version for `fastapi`, `uvicorn`).

16. `react-simple-maps` Still Installed Alongside `react-google-charts`

[package.json](#) still lists `react-simple-maps@3.0.0` even though it was previously replaced with `react-google-charts`. This adds ~150KB of dead weight to the bundle.

17. Typo in Backend Response

In [main.py:302](#):

```
return {"status": "success", ...} # Note the leading space in "status"
```

This will cause client-side parsing issues if checking `response.status`.

18. Supabase Client Uses Non-Null Assertion + Fallback Chain

In [supabase.ts](#):

```
const supabaseAnonKey = process.env.NEXT_PUBLIC_SUPABASE_ANON_KEY!
  || process.env.NEXT_PUBLIC_SUPABASE_PUBLISHABLE_DEFAULT_KEY!;
```

The `!` (non-null assertion) makes the `||` fallback dead code — TypeScript tells the compiler the first value is always truthy. Use conditional checks instead.

19. `api.ts` Makes Excessive Supabase Round-Trips

Functions like [getPatientAppointments\(\)](#) make **3 sequential Supabase queries** (appointments → doctors → profiles) that could be a single query with JOINs:

```
SELECT a.*, d.specialty, p.full_name
FROM appointments a
JOIN doctors d ON d.user_id = a.doctor_id
JOIN profiles p ON p.id = d.user_id
WHERE a.patient_id = $1
```

20. No Error Boundaries in React Components

Large page components (20K+ line dashboard files) have no React Error Boundaries. A single crashed API call takes down the entire dashboard.

● SEVERITY: IMPROVEMENT — Strategic Suggestions

21. Replace Client-Side Diagnosis Engine with Server-Side Service

The entire 1,158-line diagnosis engine runs **in the browser**. This means:

- Every user downloads 54KB of diagnostic logic
- No ability to update the model without a frontend deploy
- No audit trail of diagnoses for compliance
- Easy to reverse-engineer the scoring

Recommendation: Move to a server-side microservice (Python/FastAPI) with:

- A proper Bayesian Network engine (pgmpy, PyMC)
 - MCMC posterior sampling for calibrated confidence
 - Database-driven condition rules (updatable without deploys)
 - Audit logging for every diagnosis
-

22. Add a Proper Observability Stack

The codebase uses `console.log/error` everywhere with no:

- Structured logging (e.g., Pino, Winston)
 - Error tracking (e.g., Sentry)
 - Performance monitoring (e.g., OpenTelemetry)
 - Audit trail for medical decisions
-

23. Implement Proper TypeScript Strictness

Several files use `any` types freely:

- `api.ts` line 44: `data: any`
- `api.ts` line 13: `availability: any`
- `retrieval.ts` line 57: `as any[]`
- Multiple `.map((a: any) => ...)` casts

Enable `"strict": true` and `"noImplicitAny": true` in `tsconfig.json`.

24. Implement Database-Driven Condition Management

Currently, conditions are either:

- Hardcoded in TypeScript files (19 files in `conditions/`)
- Or fetched via vector search

A proper CMS or admin panel for managing conditions would allow medical professionals to update diagnostic rules without engineering involvement.

25. Add API Rate Limiting on Frontend Routes

The Next.js API routes (`/api/embeddings` , `/api/admin/*` , `/api/analytics`) have **no rate limiting** — only the Python backend does. These routes are directly exploitable.

26. Implement Proper Testing

The project has:

- One Playwright config but minimal E2E tests
 - Test files in the diagnosis engine (`test_clinical_engine.ts`) but no test runner or CI integration
 - No unit tests for `api.ts`, RBAC, scheduling, or any component
-

27. `package.json` Still Named `temp-app`

The package name is "temp-app" — likely from initial scaffolding. Should be "healio-ai".

Summary Priority Matrix

Priority	Issue	Impact	Effort
🔴 P0	Leaked API key in <code>.env</code>	Security breach	30 min
🔴 P0	No server-side route protection	Any user = admin	2-4 hours
🔴 P0	<code>getNotifications</code> bug	Silent data issue	5 min fix
🔴 P1	Backend has no auth	Open email spam	2-3 hours
🟡 P1	Replace pseudo-Bayesian engine	Core product quality	2-4 weeks
🟡 P2	Remove dead <code>calculateScore()</code> function	Code cleanliness	15 min
🟡 P2	Fix embedding fallback to return error, not random vectors	Diagnosis accuracy	1 hour
🟡 P2	Remove mock data from production	Bundle size, data hygiene	30 min
🟡 P3	Delete legacy dir + root dump files	Repo hygiene	1 hour
🟡 P3	Upgrade middleware cookie pattern	Auth reliability	30 min
🔵 P3	Move engine server-side	Architecture	1-2 weeks
🔵 P3	Add observability + testing	Reliability	Ongoing